

GNBG-III Competition Protocols (MATLAB Harness)

Participant Guide

Scope

This document describes the **exact evaluation protocol** implemented by the provided MATLAB harness `GNBG_III_CompetitionHarness.m` and the required participant algorithm interface. Participants should implement their method by **replacing the internal logic** inside the provided algorithm template (e.g., `runAlgorithmTemplate.m`) while keeping the interface unchanged.

Benchmark suite

The harness evaluates **24 GNBG-III problems** (files `F1_...` to `F24_...` in `GNBG_III_Benchmarks/`). Each problem loads a struct `GNBG` containing (at minimum):

- `GNBG.Dimension` (problem dimension D)
- `GNBG.MaxEvals` (evaluation budget; competition standard is **500,000**)
- `GNBG.MinCoordinate`, `GNBG.MaxCoordinate` (bounds; scalar or length- D)
- `GNBG.OptimumValue` (known optimum objective value f^*)

Special cases in `fitness.m`.

- **F23 (Noisy):** if `GNBG.NoiseLevel` exists, the returned value is perturbed by Gaussian noise.
- **F24 (Dynamic):** if `GNBG.DynamicShift` and `GNBG.DynamicPeriod` exist, the landscape is shifted at function-evaluation counts that are multiples of `DynamicPeriod` (after $\text{FE} > 0$). This happens *inside fitness*.

Run setup and randomness

For each problem F_i , the harness performs `RunNumber` **independent runs**. A deterministic seed is set per (problem, run) as:

```
rng(100000 + 1000*i + run).
```

This ensures full reproducibility as long as the harness is unchanged and your algorithm uses MATLAB's RNG.

Evaluation budget and reporting points

Each run is capped at `GNBG.MaxEvals` evaluations. The harness reports error at the following fixed evaluation points:

```
evalPoints = {10,000, 50,000, 100,000, 150,000, 200,000, 250,000, 300,000, 350,000, 400,000, 450,000, 500,000};
```

Additionally, multi-target metrics use the thresholds:

$$\text{TargetThresholds} = \{10^{-1}, 10^{-3}, 10^{-5}, 10^{-8}\}.$$

Participant algorithm interface (required)

Your algorithm is called once per run as:

```
[BestHistory, BestValue, BestPosition, GNBG, AcceptanceReachPoint] = ...
    AlgorithmHandle(GNBG, AlgorithmParams);

% Optional 6th output:
[BestHistory, BestValue, BestPosition, GNBG, AcceptanceReachPoint, Extra] = ...
    AlgorithmHandle(GNBG, AlgorithmParams);
```

Inputs

- **GNBG**: benchmark struct loaded from the `.mat` file.
- **AlgorithmParams**: a struct containing any algorithm hyperparameters you need (set in the harness).

Outputs

- **BestHistory**: **row vector** of length `GNBG.MaxEvals` storing the **best-so-far objective value** after each function evaluation (FE), i.e., `BestHistory(fe)` is the best value observed up to $FE = fe$. If you output a shorter vector, the harness pads it with the final value; if longer, it truncates.
- **BestValue**: final best objective value.
- **BestPosition**: final best decision vector (size $1 \times D$).
- **GNBG**: updated struct (notably `GNBG.FE` is advanced by `fitness` calls).
- **AcceptanceReachPoint**: FE index at which $|f_{best} - f^*| \leq 10^{-8}$ is **first reached**, or `Inf` if never reached (your template computes this from `BestHistory`).
- **Extra** (optional): struct for additional diagnostics. The harness reads: `Extra.DiversityHistory` (up to 50 values), `Extra.ImprovementCount`, `Extra.StagnationPeriods`.

How to perform evaluations (mandatory)

All evaluations **must** be performed via:

$$[f, GNBG] = \text{fitness}(X, GNBG).$$

- **X** may be a single solution ($1 \times D$) or a batch ($N \times D$).
- The function increments `GNBG.FE` internally **by 1 per evaluated solution**.
- You must **not exceed** `GNBG.MaxEvals`. When using batch evaluation, ensure that `GNBG.FE + N` does not overrun the remaining budget.

What the harness computes (per run)

Let $f_{best}(FE)$ be the best-so-far objective value at evaluation count FE, and f^* be `GNBG.OptimumValue`. The harness forms the **error curve**:

$$e(FE) = |f_{best}(FE) - f^*|.$$

BestHistory normalization

Before computing metrics, the harness:

- reshapes `BestHistory` to a row vector,
- pads/truncates it to length `MaxEvals`, and
- enforces monotonicity: if `BestHistory(fe) > BestHistory(fe-1)`, it is overwritten by `BestHistory(fe-1)`.

Therefore, even if your internal tracking is imperfect, the harness evaluates a non-increasing best-so-far trace.

Error at evaluation points

For each evaluation point $p \in \text{evalPoints}$, the harness records:

$$\text{ErrorAtPoints}(p) = e(p).$$

FE-to-target (multi-target)

For each threshold $\tau \in \{10^{-1}, 10^{-3}, 10^{-5}, 10^{-8}\}$, the harness records:

$$\text{FEto_}\tau = \min\{\text{FE} : e(\text{FE}) \leq \tau\},$$

and stores `Nan` if never reached within the budget.

Success flags

For each target τ , a run is a **success** if it reaches the target within budget (i.e., `FEto_tau` is finite). For the strictest target (10^{-8}), the harness also stores:

- `AcceptancePoints(run) = AcceptanceReachPoint` (finite) or `Inf` (failure),
- `SuccessRate` (percentage of runs with finite `AcceptancePoints`).

Expected Running Time (ERT)

For each target τ , the harness computes:

- Let FE_r be the FE-to-target for run r (`Nan` if failure).
- Replace failures by `MaxEvals`: $\widehat{\text{FE}}_r = \text{FE}_r$ if success else `MaxEvals`.
- Then

$$\text{ERT}(\tau) = \frac{\sum_{r=1}^R \widehat{\text{FE}}_r}{\#\{\text{successful runs}\}},$$

and `ERT` is set to `Inf` if there are `zero` successful runs.

Aggregated statistics (per problem)

For each problem, the harness stores per-evaluation-point summary statistics across runs:

$$\text{MeanErrors, StdErrors, MedianErrors.}$$

It also stores key end-budget summaries (notably at 500K): mean, median, std, min, and max errors, plus success rates and ERTs.

Outputs and submission artifacts

After all problems/runs finish, the harness writes:

- **MAT file:** GNBG_III_DetailedResults_<AlgorithmName>.mat
Contains Results (all raw/aggregated arrays), AlgorithmParams, evalPoints, TargetThresholds, and AlgorithmName.
- **CSV file:** GNBG_III_Detailed_Results_<AlgorithmName>.csv
Contains one row per (problem, run) with:
 - Algorithm, Problem, Run
 - Error_10K ... Error_500K
 - Acceptance_FE_1e-8, Success_1e-8
 - For each target τ : FE_to_ τ and Success_ τ

Implementation checklist (what to do / avoid)

- **Do not change** the harness logic for evaluation, targets, or statistics.
- Ensure all function evaluations go through **fitness** and **respect the budget**.
- Keep BestHistory as **best-so-far per FE**; update it *for every evaluation consumed*.
- Always return AcceptanceReachPoint as specified (or Inf if never reached).
- Boundary handling is participant-controlled (e.g., clipping, reflection, resampling), but must respect bounds.
- You may populate Extra with diagnostics; missing fields are allowed.